

Figure 1.2 — Hunchback Profile (Mean Token Representation):
ID across GPT-2 Small layers (mean of all tokens, N=500 sentences, 95% bootstrap CI)

